Mihir Bhave

Mumbai | bhavemihir7.work@gmail.com | +91 7021381113 | LinkedIn | GitHub

Experience

Software Developer, Cere Labs - Mumbai. India — Jul 2025 – Present

- Led a team of 5 to build AI-driven automation for FDA eCTD dossier generation, processing 100+ documents per dossier and significantly reducing manual effort.
- Analyzed and filtered large-scale regulatory datasets, transforming unstructured data into structured formats to improve data quality and output consistency.
- Delivered end-to-end deployments, containerizing services with Docker and releasing 4+ production deployments.
- Worked on a large-scale SAP-to-SAP and SQL-to-SQL data migration, synchronizing gigabytes of data across 50,000 tables; took ownership of a partially completed project and drove it to successful completion.
- Managed client interactions including requirement gathering, knowledge transfer, and environment setup.
- Collaborated within a 12-member Agile team, contributing to code reviews, debugging, and system design.

Volunteering

Mentor, Nakshatra: A Tech Hackathon — April 2026

- Mentored 20 teams at a national-level hackathon organized by Unstop, guiding participants in building AI/ML and LLM-based solutions.

Contributor, OpenWISP - Remote — Dec 2019 – Jan 2020

- Migrated OpenWISP-Users and OpenWISP-IPAM modules from Python 2 to Python 3 and Django 2 to Django 3.
- Resolved dependency conflicts and validated test coverage across supported environments.

Education

A.P. Shah Institute of Technology, B.E. in Computer Science and Engineering (Artificial Intelligence and Machine Learning) — Dec 2021 – June 2025

- CGPA: 7.75/10.0

Royal Junior College, HSC — June 2021

- Percentage: 80

I.E.S Chandrakant Patkar Vidyalaya, SSC — June 2019

- Percentage: 75

Projects

ReguLens – AI Compliance Copilot

- Developed an AI-driven compliance engine (Python, Streamlit) to automate gap analysis between regulatory guidelines (e.g., RBI) and internal policies.
- Implemented a two-stage retrieval pipeline (semantic search + cross-encoder reranking) to accurately match and rank regulatory clauses.
- Improved compliance scoring by leveraging cross-encoder attention to detect nuanced regulatory language (e.g., “must” vs “may”).
- Integrated Gemini 1.5 Flash to generate structured remediation plans and executive summaries, reducing manual audit effort.

Skills and Technologies

Skills: Debugging, Code Reviews, Testing
Languages & Frameworks: Python, SQL, FastAPI, React, Agile
Technologies: Vector Databases, Docker, Git, Linux, RabbitMQ, Large Language Models (LLMs)